

CS 6301: Special Topics in Computer Science - Deep Learning for NLP

Introduction

Xinya Du
Spring 2023

CS6301

- People:
 - Instructor: Xinya Du
 - Call me Xinya (Pronunciation: Shinn-ya)
 - or Dr. Du, or Prof. Du
 - TAs: TBD
 - Graders: TBD

Exciting age for NLP

- Fast progress in deep learning-based NLP systems

Generating icebreaker questions (ChatGPT)

- "If you could be any fictional character, who would you choose and why?"
- "What's a goal or dream you have that you've yet to accomplish?"
- "If you could have dinner with any historical figure, who would it be and why?"

Summarization/Generation (ChatGPT)

- ← **Tweet**



Mark Dredze
@mdredze

...

Idea for a killer app:

ChatGPT: Modify my ACL paper, which is currently over the limit by 4 lines, so that it fits into the 8 page limit.

3:13 PM · Jan 18, 2023 · **37.4K** Views

11 Retweets **1** Quote Tweet **240** Likes

Writing Entire Articles (GPT-3)



Mario Klingemann



@quasimondo

...

Another attempt at a longer piece. An imaginary Jerome K. Jerome writes about Twitter. All I seeded was the title, the author's name and the first "It", the rest is done by #gpt3

Here is the full-length version as a PDF:

drive.google.com/file/d/1qtPa1c...

The importance of being on twitter

by Jerome K. Jerome
London, Summer 1897

It is a curious fact that the last remaining form of social life in which the people of London are still interested is Twitter. I was struck with this curious fact when I went on one of my periodical holidays to the sea-side, and found the whole place twittering like a starling-cage. I called it an anomaly, and it is.

I spoke to the sexton, whose cottage, like all sexton's cottages, is full of antiquities and interesting relics of former centuries. I said to him, "My dear sexton, what does all this twittering mean?" And he replied, "Why, sir, of course it means Twitter." "Ah!" I said, "I know about that. But what is Twitter?"

"It is a system of short and pithy sentences string together in groups, for the purpose of conveying useful information to the initiated, and entertainment and the exercise of wits to the initiated, and entertainment and the exercise of wits to the rest of us."

"Very interesting," I said. "Has it a name?"

"It has," he said; "it is called Twitter."

"Yes," I said, "I know that, but what is it?"

"It is a system of information," he said.

"Oh, yes," I replied; "but what is it?"

"Why, sir," he said, "you can go up to any of the gentlemen you see twittering in the

https://twitter.com/quasimondo/status/1284509525500989445?ref_src=twsrc

Even for Writing Code... (GPT-3)



Sharif Shameem
@sharifshameem

...

This is mind blowing.

With GPT-3, I built a layout generator where you just describe any layout you want, and it generates the JSX code for you.

W H A T

Describe a layout.

Just describe any layout you want, and it'll try to render below!

large text in red that says "WELCOME TO MY NEWSLETTER" and a blue button

Generate

```
<h1 style={{color: 'red', fontSize: '100px'}}>WELCOME TO MY NEWSLETTER</h1><div style={{color: 'blue', padding: 20}}>Subscribe</div>
```

WELCOME
TO MY
NEWSLETTER

1:34 | 1.8M views

https://twitter.com/sharifshameem/status/1282676454690451457?ref_src=twsrc

Reasoning (GPT-3)

(b) Few-shot-CoT

Q: Roger has 5 tennis balls. He buys 2 more cans of tennis balls. Each can has 3 tennis balls. How many tennis balls does he have now?

A: Roger started with 5 balls. 2 cans of 3 tennis balls each is 6 tennis balls. $5 + 6 = 11$. The answer is 11.

Q: A juggler can juggle 16 balls. Half of the balls are golf balls, and half of the golf balls are blue. How many blue golf balls are there?

A:

(Output) *The juggler can juggle 16 balls. Half of the balls are golf balls. So there are $16 / 2 = 8$ golf balls. Half of the golf balls are blue. So there are $8 / 2 = 4$ blue golf balls. The answer is 4.* ✓

Reasoning (GPT-3)

(d) Zero-shot-CoT (Ours)

Q: A juggler can juggle 16 balls. Half of the balls are golf balls, and half of the golf balls are blue. How many blue golf balls are there?

A: **Let's think step by step.**

(Output) There are 16 balls in total. Half of the balls are golf balls. That means that there are 8 golf balls. Half of the golf balls are blue. That means that there are 4 blue golf balls. ✓

Description → Images (multimodal)

- DALL-E Generation

“Photo of hip hop cow in a denim jacket recording a hit single in the studio”



Let's go back to natural language

- deep learning a little later ...

What is the goal of human language?

What is the goal of human language?

To share thoughts with conspecifics

- to **communicate** with others
- to store and pass **knowledge**

What is the goal of human language?

To share thoughts with conspecifics

- to **communicate** with others
- to store and pass **knowledge**

To think more complex thoughts

- **reasoning**
- **interpretation**

What is the goal of human language?

- To store and pass knowledge (Extraction to construct Graph)




Figure 2: COVID-KG Overview: From Data to Semantics to Knowledge

What is the goal of human language?

- To store **knowledge** (In the form of model “parameters”)




Figure 1: A diagram of our text-to-text framework. Every task we consider—including translation, question answering, and classification—is cast as feeding our model text as input and training it to generate some target text. This allows us to use the

What is the goal of human language?

To think more complex thoughts

- **reasoning (commonsense)**





Figure 1: A tiny subset of ATOMIC, an atlas of machine commonsense for everyday events, causes, and effects.


What is the goal of human language?

To think more complex thoughts

- **reasoning (deductive)**
- **interpretation**



(a) Formal language as knowledge representation and symbolic reasoner



(b) Natural language as knowledge representation and PLM as reasoner

Figure 1: Comparison between the previous paradigm which uses formal representation and symbolic reasoner, and the new paradigm which uses natural language as knowledge representation and PLM as reasoner.

What is the goal of human language?

To think more complex thoughts

- **reasoning (abductive)**
- **interpretation**

Short fact 1	Short fact 2	Short fact 3	Rule
<p>The Venus flytrap is a carnivorous plant native to subtropical wetlands on the East Coast of the United States in North Carolina and South Carolina. It catches its prey—chiefly insects and arachnids—with a trapping structure formed by the terminal portion of each of the plant's leaves, which is triggered by tiny hairs on their inner surfaces.</p>	<p>Pitcher plants are several different carnivorous plants which have modified leaves known as pitfall traps—a prey-trapping mechanism featuring a deep cavity filled with digestive liquid. The traps of what are considered to be "true" pitcher plants are formed by specialized leaves. The plants attract and drown their prey with nectar.</p>	<p>Drosera, which is commonly known as the sundews, is one of the largest genera of carnivorous plants, with at least 194 species. The trapping and digestion mechanism of Drosera usually employs two types of glands: stalked glands that secrete sweet mucilage to attract and ensnare insects and enzymes to digest them, and sessile glands that absorb the resulting nutrient soup.</p>	<p>If a plant is carnivorous, then it probably has a trapping structure.</p>

Table 1: An example of inductive reasoning in DEER dataset. We embolden the words in facts that contain the key information to induce this rule (just to explain the relation between facts and rule, in DEER there's no special word annotations for fact).

What is the goal of human language?

To share thoughts with conspecifics

- to **communicate** with others
- to store and pass **knowledge**

To think more complex thoughts

- **reasoning**
- **interpretation**



Almost all of your use of language is internal. Virtually all of the use of language has nothing to do with communication. The idea that language has evolved as a system of communication, or designed for communication, makes no sense.

Noam
Chomsky

...

The systems of thought use linguistic expressions for reasoning, interpretation, organizing action, and other mental acts.

What is the goal of human language?

To share thoughts with conspecifics

- to **communicate** with others
- to store and pass **knowledge**

To think more complex thoughts

- **reasoning**
- **interpretation**



Almost all of your use of language is internal. Virtually all of the use of language has nothing to do with communication. The idea that language has evolved as a system of communication, or designed for communication, makes no sense.

Noam
Chomsky

...

The systems of thought use linguistic expressions for reasoning, interpretation, organizing action, and other mental acts.



The limits of my language mean the limits of my world.

Ludwig
Wittgenstein

Wouldn't be nice if computers could understand human language?



The Gutenberg Bible at Yale Beinecke Rare Book Library

Natural Language Processing (NLP)

It deals with **natural languages** that humans speak (e.g., English, Chinese, Spanish, Arabic, etc), not programming languages.

- NLP is also referred to as **Human Language Technology**

Natural Language Processing (NLP)

It deals with **natural languages** that human speak (e.g., English, Chinese, Spanish, Arabic, etc), not programming languages.

- NLP is also referred to as **Human Language Technology**

It aims to use computational approaches to analyze human languages.

- morphology, lexical analysis, syntax, semantics, discourse

Natural Language Processing (NLP)

It deals with **natural languages** that humans speak (e.g., English, Chinese, Spanish, Arabic, etc), not programming languages.

- NLP is also referred to as **Human Language Technology**

It aims to use computational approaches to analyze human languages.

- morphology, lexical analysis, syntax, semantics, discourse

It aims to build computer systems that can **understand** and **generate** natural language.

- It can be used for human-human communication (e.g., machine translation, grammatical error correction) or human-machine communication (e.g., dialog systems, question answering).

Natural Language Processing vs Computational Linguistics

A lot of people use them interchangeably.

They are very similar, but also different.

The focus is a bit different:

- Computational Linguistics focus on scientific questions in "Linguistics". It aims to study scientific questions about human languages using computational tools.
- Natural Language Processing focus on engineering problems of building practical systems and applications that can mimic human usage of languages.

Machine Translation

DETECT LANGUAGE

CHINESE

ENGLISH

SPANISH



CHINESE (SIMPLIFIED)

ENGLISH

SPANISH



到目前为止，还没有生物标志物可以预测个体对COVID-19的敏感性。
南方科技大学、上海交大、武汉中南医院、武汉金银潭医院等8家单位
的最新研究显示：A、B、O、AB血型与新冠肺炎易感性存在关联。这
是该领域的首份研究。

Dào mùqian wéizhǐ, hái méiyǒu shēngwù biāozhì wù kěyǐ yùcè gètǐ duì COVID-19 de
mǐngǎnxìng. Nánfāng kējì dàxué, shànghǎi jiāodà, wǔhàn zhōngnán yīyuàn, wǔhàn
jīnyíntán yīyuàn děng 8 jiā dānwèi de zuìxīn yánjiū xiǎnshì: A, B, O, AB xiěxíng yǔ xīnguān

Show more



107/5000

拼



To date, there are no biomarkers that can predict an individual's sensitivity to COVID-19. The latest research from 8 units including Southern University of Science and Technology, Shanghai Jiaotong University, Wuhan Zhongnan Hospital, Wuhan Jinyintan Hospital, etc. showed that blood types of A, B, O, and AB are associated with susceptibility to new coronary pneumonia. This is the first study in this field.



Question Answering (Information Retrieval)

What is the mascot for the Penn State football team

All News Images Maps Shopping More Settings Tools

About 3,340,000 results (0.86 seconds)

Nittany Lion

The Nittany Lion is the mascot of the Penn State **Nittany Lions**—the athletic teams of the Pennsylvania State University, located in University Park, Pennsylvania, USA. It is an eastern mountain **lion**, the "Nittany" forename referring to the local Mount **Nittany**, which overlooks the university.



en.wikipedia.org › wiki › Nittany_Lion ▾

[Nittany Lion - Wikipedia](#)

Question Answering (Span Extraction)

What is the mascot for the Penn State football team

All News Images Maps Shopping More Settings Tools

About 3,340,000 results (0.86 seconds)

Nittany Lion

The **Nittany Lion** is the mascot of the Penn State **Nittany Lions**—the athletic teams of the Pennsylvania State University, located in University Park, Pennsylvania, USA. It is an eastern mountain **lion**, the "Nittany" forename referring to the local Mount **Nittany**, which overlooks the university.

[en.wikipedia.org › wiki › Nittany_Lion](https://en.wikipedia.org/wiki/Nittany_Lion)


[Nittany Lion - Wikipedia](#)




Conversational Personal Assistants



Apple Siri




Google Assistant



Microsoft Cortana



Ambiguity

- Linguistic perspective ...





Natural Language Understanding is Hard

"The old man the boat."



Natural Language Understanding is Hard

"The old man the boat." syntax ambiguity



Semantic Ambiguity

At last, a computer that understands you like your mother.

Semantic Ambiguity

At last, a computer that understands you like your mother.

- Direct meanings:
 - It understands you like your mother (does) [presumably well]
 - It understands (that) you like your mother

Semantic Ambiguity

At last, a computer that understands you like your mother.

- Direct meanings:
 - It understands you like your mother (does) [presumably well]
 - It understands (that) you like your mother

Context!

- Thus, the **key difference** between semantics and pragmatics is the fact that semantics is context independent whereas **pragmatic is context dependent**.

Context!



1. Stick your cigarette butt here
2. Stick your backside here

what type of context?

Ambiguities in the Wild: Context




1. Stick your cigarette butt here (trash can)
2. Stick your backside here

Even extra-linguistic features



Q: How many home runs has Shohei Ohtani hit?

A: 24



Q: How many home runs has Shohei Ohtani hit?


A: **25**



Even extra-linguistic features


Context Type: Temporal

Question: Which COVID-19 vaccines have been authorized for adults in the US?



Context Type: Geographical

Question: Which COVID-19 vaccine was the first to be authorized by our government?



Even extra-linguistic features

- Time
- Location
- Demographics
- Culture
- ...
- ? assuming we have a strong Neural Language Model that can accurately answer questions (e.g. ChatGPT). How do we take into account extra-linguistic, keeping updated to the “current dynamic world”?

NLP Tasks

- [Automatic speech recognition](#)
- [CCG](#)
- [Common sense](#)
- [Constituency parsing](#)
- [Coreference resolution](#)
- [Data-to-Text Generation](#)
- [Dependency parsing](#)
- [Dialogue](#)
- [Domain adaptation](#)
- [Entity linking](#)
- [Grammatical error correction](#)
- [Information extraction](#)
- [Intent Detection and Slot Filling](#)
- [Language modeling](#)
- [Lexical normalization](#)
- [Machine translation](#)
- [Missing elements](#)
- [Multi-task learning](#)
- [Multi-modal](#)
- [Named entity recognition](#)
- [Natural language inference](#)
- [Part-of-speech tagging](#)
- [Paraphrase Generation](#)
- [Question answering](#)
- [Relation prediction](#)
- [Relationship extraction](#)
- [Semantic textual similarity](#)
- [Semantic parsing](#)
- [Semantic role labeling](#)
- [Sentiment analysis](#)
- [Shallow syntax](#)
- [Simplification](#)
- [Stance detection](#)
- [Summarization](#)
- [Taxonomy learning](#)
- [Temporal processing](#)
- [Text classification](#)
- [Word sense disambiguation](#)

NLP Tasks

Foundational Technologies

- Language Modeling
- Part-of-speech Tagging
- Syntactic Parsing
- Dependency Parsing
- Named Entity recognition
- Coreference resolution
- Word Sense Disambiguation
- Semantic Role Labelling
-

High-Level Tasks and Applications

- Sentiment Analysis
- Information Extraction
- Machine Translation
- Question Answering
- Semantic Parsing
- Summarization
- Dialogue systems
- Language and Vision
- Data-to-Text Generation
-

Natural Language Processing Systems

Input X	Output Y	Task
Text (e.g., Sentiment Analysis)	Label	Text Classification
Text Tagging)	Linguistic Structure	Structured Prediction (e.g., Part-of-Speech
Text (e.g., Translation)	Text	Text Generation

Natural Language Processing Systems

To create systems for X → Y, we can

- **Expert Systems:** Manually create rules mapping X to Y based on our knowledge of languages (below is an example for detecting named entities)

PER, POSITION of ORG:

George Marshall, Secretary of State of the United States

PER (named|appointed|chose|etc.) PER Prep? POSITION

Truman appointed Marshall Secretary of State

PER [be]? (named|appointed|etc.) Prep? ORG POSITION

George Marshall was named US Secretary of State

Hand-built patterns have the advantage of high-precision and they can be tailored to specific domains. On the other hand, they are often low-recall, and it's a lot of work to create them for all possible patterns.

- **Data-Driven Methods:** Collect Data and Use Machine Learning Techniques to learn the mapping function


Data Driven Methods and Machine Learning

“**Machine Learning**” is a data-driven problem solving strategy.

Data Driven Methods and Machine Learning

“Machine Learning” is a data-driven problem solving strategy.

The ultimate goal of machine learning is **generalization!** Making future prediction on unseen examples as accurate as possible!




Common theme is to solve a prediction problem:


Given an input \mathbf{x} ,

Predict an “appropriate” output \mathbf{y} (label, decision, action, etc).

Machine Learning Pipeline



Training, Validation, Test



Data Hygiene

Training Data: Use this and only this to train your neural networks.

Development Data: Use this to decide when to stop training your neural networks, and do hyperparameter tuning.

Test Data: Use test data after picking your model. You should use test data only once.

Neural Networks == Deep Learning






Figure 1: The Transformer - model architecture.

Deep Learning for NLP



From “Real-Time Social Media Analytics with Deep Transformer Language Models: A Big Data Approach”
by Ahmet and Abdullah

Keep updated with NLP research

- <https://aclanthology.org>

Keep updated with NLP research

- Twitter

The image shows a screenshot of a Twitter feed. The first tweet is from **William Wang** (@WilliamWangNLP) dated Jan 18. It discusses in-context learning / few-shot prompting at inference time, comparing it to an #NLProc version of explicit self-supervised imitation learning (SIL) of test-time adaption from [@xwang_lk](#)'s 2019 CVPR paper. The tweet includes a link to arxiv.org. The second tweet is from **Yonatan Belinkov** (@boknilev) dated Jan 18, responding to the first tweet with an interesting read and a link to arxiv.org/abs/2301.05272.

William Wang @WilliamWangNLP · Jan 18

IMHO, in-context learning / few-shot prompting at the inference time can be seen as an [#NLProc](#) version of explicit self-supervised imitation learning (SIL) of test-time adaption from [@xwang_lk](#)'s 2019 CVPR paper.

arXiv arxiv.org
Reinforced Cross-Modal Matching and Self-Supervi...
Vision-language navigation (VLN) is the task of navigating an embodied agent to carry out natural ...


Like 3 1,597

Yonatan Belinkov @boknilev · Jan 18


Interesting read, although I doubt the closed nature of LLMs would lead to reintroduction of rule based approaches.
A few other interesting thoughts as well.
[arxiv.org/abs/2301.05272](#)

Keep updated with NLP research

- Twitter



A screenshot of a Twitter profile for a user named AK (@_akhaliq). The profile picture shows a landscape with a small house and mountains at night. The bio reads: "paper tweets, dms are open, ML @Gradio (acq. by @HuggingFace 😊) instagram: [instagram.com/_akhaliq/](https://www.instagram.com/_akhaliq/) discord: discord.gg/Vef7GyFVS8". Below the bio, it says "Science & Technology" and provides a LinkedIn link (linkedin.com/in/ahsenkhaliq). The account was joined in April 2014. The following statistics are listed: 1,691 Following and 113.3K Followers. A note below states: "Followed by Clara Na, Karsten Kreis, and 74 others you follow". At the bottom, there are tabs for "Tweets", "Tweets & replies", "Media", and "Likes", with "Tweets" currently selected. A pinned tweet from December 18, 2022, says: "I will start posting arxiv paper tweets on Patreon (public)".



A screenshot of a Twitter profile for the account "arXiv cs.CL Papers" (@arxiv_cs_cl). The profile picture is red with the text "arXiv cs.CL". The bio reads: "Unofficial account for updates of Computation and Language on arXiv". The account was joined in December 2015 and has 2 Following and 4,640 Followers. It is followed by Jialu Li, Zonglin Yang, and 46 others. Below the bio, there are tabs for "Tweets", "Tweets & replies", "Media", and "Likes", with "Tweets" selected. A recent tweet from the account is shown: "cs.CL Papers @arxiv_cs_cl · 4h ift.tt/XcldNTv Understanding and Detecting Hallucinations in Neural Machine Translation via Model Introspection. (arXiv:2301.07779v1 [cs.CL]) #NLProc". To the right, there is a sidebar with a logo for "cs.cl" and the text: "Inicio - Covarrubias & Cía. Covarrubias & Cía. es una firma especializada y enfocada exclusivamente en la propiedad intelectu...".

Course Topics

- Part 1: Introduction and Neural Network Basics
 - Introduction to NLP (this lecture)
 - Text Classification and Language Modeling
 - Neural Networks and Backpropagation

Course Topics

- Part 1: Introduction and Neural Network Basics
 - Introduction to NLP (this lecture)
 - Text Classification and Language Modeling
 - Neural Networks and Backpropagation
- Part 2: DL&NLP Models, Applications
 - Sequence-to-Sequence, Attention, Transformers
 - Pretraining Language Models and Transfer Learning
 - Text-to-Text and Prompt-based Learning

Course Topics

- Part 1: Introduction and Neural Network Basics
 - Introduction to NLP (this lecture)
 - Text Classification and Language Modeling
 - Neural Networks and Backpropagation
- Part 2: DL&NLP Models, Applications
 - Sequence-to-Sequence, Attention, Transformers
 - Pretraining Language Models and Transfer Learning
 - Text-to-Text and Prompt-based Learning
- Part 3: Advanced Techniques (with paper discussions)
 - Retrieval-augmented models
 - In-context learning
 - Prompting
 - Vision-language Models
 - ..

Logistics

Course Information

Lecture

- Friday, 10:00am - 12:45pm (11:15--11:30 resting period)
- GR 2.530, In class attendance is encouraged.

Instructor

- Xinya Du, Office Hour (in person): ECSS 3.227, Friday 2 – 3pm
- TA, Office Hour: TBD

How to Contact us

- For most technical questions
 - ask on eLearning “Q&A” forum
- Please **email** TA.
- Do not send DMs via Teams

The screenshot shows a 'Discussion Board' section with a sub-section titled 'Create Forum'. Below this, there is a list of forums. One forum is visible, titled 'Q&A', with a description: 'Discussions regarding lectures and paper'. There are 'Delete' buttons next to each forum entry.

FORUM	DESCRIPTION
Q&A	Discussions regarding lectures and paper

Textbooks

Available online! You are encouraged to read beyond slides.

-  [Neural Network Methods in Natural Language Processing](#) Yoav Goldberg.
 - <https://link.springer.com>
-  [Speech and Language Processing](#). 3rd edition. Dan Jurafsky and James H. Martin.
- [Natural Language Processing](#). Jacob Eisenstein.

Prerequisites

- Programming in Python
- CS 5343 Algorithm Analysis and Data Structures
- CS 6375 Machine Learning
- CS 6320: Natural Language Processing (not a must)

Previous experience in NLP/Deep Learning is preferred but not required.

Groups

- Paper presentation and the final project are group-based.
- All group members are treated equally and earn the same grades.
- Fill in the Excel Sheet (link will be posted on eLearning) with group member's names and NetIDs.

	A	B	C	D	E
1	please form groups and fill in your names in the cells below, each group should consists of <=4 students, you can use the discussion board to look for teammates. The ddl for grouping is Jan 25.				
2		member 1	member 2	member 3	member 4
3	group0	Xinya Du (xxd220001)	TA		
4	group1				
5	group2				

Groups

- Each group should contain **2-4** students. If you detect serious imbalance workload among group members, penalty will apply
- Utilizing the forum to find group members
- Deadline for grouping is Jan 25 <cs6401-spcs6401-sp23-groups.xlsx>

	A	B	C	D	E
1	please form groups and fill in your names in the cells below, each group should consists of <=4 students, you can use the discussion board to look for teammates. The ddl for grouping is Jan 25.				
2		member 1	member 2	member 3	member 4
3	group0	Xinya Du (xxd220001)	TA		
4	group1				
5	group2				

Groups

- Link will be posted on eLearning: [cs6401-sp23-groups.xlsx](#)

	A	B	C	D	E
1	please form groups and fill in your names in the cells below, each group should consists of <=4 students, you can use the discussion board to look for teammates. The ddl for grouping is Jan 25.				
2		member 1	member 2	member 3	member 4
3	group0	Xinya Du (xxd220001)	TA		
4	group1				
5	group2				

Grading

- Assignments (30%)
- Paper presentation (20%)
- Course project (45%)
- Participation (5%)

Assignments (individual-based)

Two Assignments

- Assignment 1 (Written)
 - Time: TBD
 - Related to basic concepts (e.g. word representations, backprop, FFNN)
- Assignment 2 (programming)
 - Time: TBD
 - Related to **correct** implementation of a neural network model for classification task

Not late assignment accepted unless there is a compelling reason. Late assignment: please refer to the syllabus.

Paper Presentation (group-based)

Research paper

- We will provide potential topics (each with 1-2 papers, you can also add in related papers!)

Paper Presentation (group-based)

Research paper

- We will provide potential topics (each with 1-2 papers, you can also add in related papers!)
- **Paper Review (5%)**
 - Send me the slides (cc'ing the TA) 48 hours before your group's presentation
 - After the presentation, upload the slides to eLearning.

Paper Presentation (group-based)

Research paper

- We will provide potential topics (each with 1-2 papers, you can also add in related papers!)
- **Paper Review (5%)**
 - Send me the slides (cc'ing the TA) 48 hours before your group's presentation
 - After the presentation, upload the slides to eLearning.
- **Paper presentation (15%)**
 - Time: in class, 35 minutes for each group
 - We value clear presentations and insightful discussions
 - You are allowed to search materials/explanations online, but you must need to understand the contents and demonstrate your understanding of this topic.

Exams

No Exam. This course is a **research project** oriented course.

Project Format

Project Topics. This project aims to conduct original and independent research over **any NLP-related topics**.

- Must deal with natural language data
- Must include some degree of model implementations and experiments

Group Policy. You can work on the course project in a group of **2 - 4** people. You are allowed to combine this project with your research projects or projects from other courses.

Project Deliverables

- **Proposal.** Write a proposal that outlines your plan including what problem or task you want to address, what dataset(s) you want to work on, what metrics you need to employ, what baselines you would like to compare with. You should also cite a few relevant prior papers. The length should be 2—4 pages.
- **Proposal Oral Presentation.** During the mid of the semester, each group will be given a 15—20 min oral presentation.

Project Deliverables

- **Final Report.** Your final report should use the ACL Latex template (<https://acl-org.github.io/ACLPUB/formatting.html>), overleaf template [here](#). The length should be around 4—8 page plus references. Your report should begin with an abstract and introduction to clearly state the problem you want to solve and contributions you have made. It should also have a section on related work, a section on your methodology, a section on your experimental settings and results, and a section on conclusions.
- **Final Oral Presentation.** During the last week (and the week before) of the lecture, each group will give a 15—20 min oral presentation.

Participation

- Quiz (2%)
- Discussions (2%)
 - Asking and answering other students' questions on the "Q&A" forum
 - Asking and participating in other groups' presentations
 - Including (1) elearning discussion forum; (2) class question answering and asking. Top 30% students get 2%, 30%-70% students get 1%.
- Course feedback (1%)

Participation

- Quiz (2%)
 - 1-2 quiz, each contains multi-choice questions or short-form written questions.
 - within 10 minutes

Tentative schedule

Week	Date	Topic	Assignments	Candidate Papers	Recommended Readings
Week 1	Jan 20	Introduction -			
Week 2	Jan 27	Neural Networks, Backpropagation -	a1		J&M 7 (.1--.4), Primer, J&M 7, Intro to Computation Graphs
Week 3	Feb 3	Word Embeddings -			J&M 6, word2vec explained
Week 4	Feb 10	Pytorch & Transformers Tutorial RNNs, Seq2Seq, Attention	a1 due, a2		J&M 9, J&M 10 (.2, .3), Luong 15
Week 5	Feb 17	- Self-attention			illustrated Transformer , Annotated Transformer, Paper
Week 6	Feb 24	- Pretrained Language Models (PLMs)	a2 due		Survey
Week 7	Mar 3	- DL for NLP applications (QA, NLG)			
Week 8	Mar 10	Project proposal Project proposal		project proposal report due	
Week 9	Mar 17	Spring break			

Tentative schedule

Week 9	Mar 17	Spring break			
Week 10	Mar 24	In-context learning Prompting			Survey
Week 11	Mar 31	Paper Presentation * 3 Paper Presentation * 3			
Week 12	Apr 7	Retrieval-augmented Models -			
Week 13	Apr 14	Paper Presentation * 3 Paper Presentation * 3			
Week 14	Apr 21	Vision-Language Models -			
Week 15	Apr 28	Paper Presentation * 3 Paper Presentation * 3			
Week 16	May 5	Project final presentation Project final presentation	project final report due		

End of Introduction